



Illustration: David Evans/Bloomberg Law

Bloomberg Law

Artificial Intelligence: The Next Phase

As AI moves from a question of “what” to “how,” in-house counsel wrestle with issues around liability, data governance, and compliance.

Table of Contents

EXPLAINER

- 3** Five Artificial Intelligence Questions for In-House Legal Teams
1. How do I comply with AI laws?
 2. Should I commit to voluntary standards?
 3. How do I vet my vendors?
 4. How do I get my board on board?
 5. How can AI make my job easier?

FEATURE STORY

- 8** Assigning Blame for AI Harms Raises Tricky Liability Questions

PRACTITIONER INSIGHT

- 10** GCs Can Set Tone, Serve as Guides as Companies Navigate AI Era

PRACTITIONER INSIGHT

- 12** Open Dialogue Sets Up In-House Team for AI Adoption Success

PRACTITIONER INSIGHT

- 14** How In-House Legal Departments Should Engage with Generative AI

PRACTITIONER INSIGHT

- 16** E-Discovery's Next Frontier Is Your AI Tool's Output, Prompt Log

- 18** Glossary of Terms

EXPLAINER

Five Artificial Intelligence Questions for In-House Legal Teams

By Isabel Gottlieb

Questions about corporate governance of AI have moved quickly from the hypothetical to the "how."

Companies are seeking to determine appropriate standards and guardrails for using the technology responsibly, and also who should direct those governance efforts.

Nearly two years after ChatGPT kicked off the frenzy, the challenges have become much clearer, and it's up to companies to meet them.

"A lot of people are trying to figure out how to go from principles to practice," said Evi Fuelle, global policy director at Credo AI, an AI governance platform.

An array of nascent laws—from the EU to a smattering of U.S. states and elsewhere in the world—makes AI governance increasingly a compliance question. Even where there's no regulation, voluntary standards can get companies on track to comply with future laws.

Mass uptake of AI also gives general counsel and chief legal officers a reason to examine their own roles, perhaps in moving "from blocker to wise enabler," said Sebastian Niles, the chief legal officer at Salesforce and a former Wachtell Lipton partner.

"You need to figure out where, maybe, there are things to block," he added.

But the more important issues are, "Where and how can that CLO facilitate the understanding of: What should we be enabling? What should we be doing?"

General counsel and C-suite executives charged with navigating through AI's risks

face additional challenges, including how to get buy-in from corporate boards, how to evaluate third-party vendors, and how to effectively use AI within their own legal departments.

1 How do I comply with AI laws?

For many companies, step one is figuring out how to comply with the EU's sweeping Artificial Intelligence Act—the most fully-formed AI law in force today. Provisions will take effect in stages, with many of the key rules starting to apply in August 2026. Rules barring unacceptably risky AI, which also carry the most punitive fines, go live in February 2025.

"The adoption of the EU AI Act really puts its thumb down," said Lee Tiedrich, a professor at Duke University and an adviser to the OECD's AI work. "We now have comprehensive AI legislation in the EU, which is going to have effects beyond the EU. That's a significant change."

Many companies are applying the EU requirements for high-risk systems across the board, said Philip Dawson, head of global AI policy at Armilla AI, which does AI testing to help companies comply with laws and standards.

"They may want to have one process, across their company," he said. Broadly applying stricter rules simplifies the tracking of AI system risk classifications, which can change according to how the systems are used.

The U.S. lags Europe at the federal level. In the absence of action from Congress, at least 18 U.S. states have enacted laws addressing issues such as election

EU AI Act Timeline

- August 2, 2024
EU AI Act into force
- February 2, 2025
Prohibition on unacceptably risky AI in effect
- August 2, 2025
General purpose AI rules in effect
- August 2, 2026
Most high-risk AI rules in effect
- August 2, 2027
Remaining provisions in effect

Visit Bloomberg Law's In Focus: Artificial Intelligence center for a full list of resources on the EU AI Act.

[Click to learn more.](#)

deepfakes, phony intimate images, and discrimination by automated systems. Only Colorado has adopted a comprehensive measure aimed at regulating how the private sector uses the technology.

Regulatory approaches likely won't align perfectly across national and state lines, Fuelle noted, but corporate AI leaders are trying to identify common ground across jurisdictions.

She said there are recurring themes among those approaches, including requirements for: inventorying AI uses, testing, rooting out security flaws, the need for human oversight, independent evaluation, and ongoing risk management.

U.S. AI regulations are aimed mostly at having companies test systems for fairness and bias. Such measures include insurance regulations in Colorado and New York, and a New York City law governing AI employment tools.

But while regulators are setting expectations for how companies should test their systems for dangers like racial bias, Dawson said, "The compliance tools or technical solutions have not yet been approved. Translating the policy or regulatory requirements into technical solutions—that is being worked out."

Many of his clients have charged internal data science teams with testing proprietary and third-party AI applications for safety, as well as for bias risks, he said. But that work becomes ever more challenging as companies increasingly experiment with generative AI applications and AI tools that are frequently updated.

"So there's a bit of a backlog that's either exposing them to risk and liability or compliance issues down the road," he added, and as a result, many are turning to specialized third-party testing companies.

A company may approach regulatory compliance differently if it's using third-party tools or developing its own,

said Shelley McKinley, the chief legal officer at GitHub.

"Using AI tools is often just a matter of having a sensible, easy-to-understand internal use policy," she said. Enough companies have adopted them by now that they can provide a good model, she added—and organizations like GitHub offer customers a sample policy.

But for companies developing AI tools, compliance may get more involved as the regulatory landscape quickly evolves, McKinley said.

She recommended investing in subject matter experts to track regulatory requirements, adapting existing policies concerning data governance, and understanding the elements of AI products.

"Most AI solutions will be a blend of open source and proprietary models, code, and platforms," McKinley said. She recommended asking third-party providers for clear documentation on topics like data privacy, and asking providers for solutions that can help when AI is deployed in higher-risk scenarios, "where the transparency and compliance obligations may be higher, and where dedicated tools and filters are critical to meeting AI regulatory requirements."

Stay Up to Date State Ethics



Track state legal ethics guidance related to AI and professional responsibility.

2 Should I commit to voluntary standards?

As laws and regulations accumulate, some companies are looking to voluntary standards and certifications to guide them before mandatory rules come into play.

Organizations are looking for, "What's the stamp that we can get as an organization to say, 'Hey, we do AI governance right?'" said Ryan Donnelly, co-founder of the AI governance platform Enzai.

Donnelly said he is often asked about the International Organization for Standardization's ISO 42001, which outlines how to set up internal policies and procedures to best manage AI risks. Companies that want to be seen as first movers in the AI governance space are looking toward 42001, he said.

In the U.S., the National Institute of Standards and Technology, which sits under the Commerce Department, has released an AI Risk Management Framework focused on governance.

ISO certification could be valuable if the biggest customers of the big AI companies begin demanding it. The NIST framework, a guide to making good decisions about AI rather than a certification, can sit alongside the ISO standard, said Craig Shank, a consultant on AI ethics and formerly VP of corporate, external and legal affairs at Microsoft.

Even if a company doesn't obtain the full ISO certification, it's "a useful grounding point to say, 'we're making deliberate decisions,'" including documenting how a decision about AI was made, Shank said.

"I think it is as protective as just about anything you can do from a liability perspective, if for no other reason than it makes it much faster to unwind a problem if you've created a mess," he added. "If you know the decision points that you made in order to get it there, you at least know where the off button is, or you know where to unwind."

3

How do I vet my vendors?

At GitHub, vetting AI tools starts with the same questions the company would ask about any other tool, McKinley said, looking at issues like, "vendor reputations, data handling, security and privacy protections."

"In a lot of respects we approach these tools like we've long approached any third-party tools we use: Know

your vendor, their reputation, and the assurances and transparency they provide about their offerings," she added.

Josh Nicosia, chief legal officer at Smoothie King, runs a small legal department—just one other attorney and two paralegals. When he went searching for a contract lifecycle management tool to help his department run more efficiently, he said, his first questions for vendors were, "Is our information shared? What are you using it for?"

By now, though, most major legal AI vendors say they won't train on clients' information, he said.

The company is also adding provisions to its contracts with all third-party vendors—not just legal tools—that they must disclose when any deliverable is created using AI, and what platform it used. Smoothie King's contracts also have clauses that allow it to terminate any for which it deems AI use to be inappropriate.

One big concern is a vendor using AI to generate marketing materials—and inadvertently violating someone else's copyright or trademark, leading to lawsuits and reputational damage, Nicosia said.

Contracts are an important tool to protect companies from legal risks involving third-party AI vendors, but shouldn't be the only one.

"If contracting around risk is your only method of risk mitigation, you're in trouble," said Devika Kornbacher, co-chair of the global tech group at Clifford Chance.

One of the most important questions companies can ask their AI vendors today is how they're keeping up with the constant emergence of new models, said Arun Subramaniyan, founder and CEO of Articul8, a company that builds domain-specific models for clients.

Generative AI tools are built on one or more underlying models. No matter how good those models are, something better

"If contracting around risk is your only method of risk mitigation, you're in trouble."

will come along quickly, Subramaniyan said. But a vendor shouldn't have to completely rebuild each time they adjust to include a new model, he said. And while companies should be diligent about testing tools for risks like bias, he said, if the vendor has done their work well, the diligence process shouldn't need to restart from scratch.

Vendors should also be transparent about how they're testing their own systems, including what metrics they're using, he said.

He also recommended asking potential vendors to be up front about the actual cost of the AI tool.

"Your total cost of ownership is completely hidden," he said. For example, the vendor may quote a price based on tokens, but that doesn't take into account whether the client needs an internal compliance team to use the tool, or whether they can use the vendor's.

"When you sell enterprise software, usually the enterprise software vendor has to attest for the compliance of the software they put in," he said. "These sets of products are so new, nobody's even asking those questions."

4 How do I get my board on board?

How general counsel talk to their boards depends in part on how much AI knowledge members have, said Rob Chesnut, Bloomberg Law columnist and former general counsel and chief ethics officer at AirBNB.

AI has only recently become a concern for most companies, which means many corporate boards won't have a member who's an AI expert, he said.

"For me, as a GC, I'm looking at my matrix," or how the roster of corporate board members aligns with key areas, Chesnut said. "I'm deciding whether

AI is, in fact, a fundamental quality of expertise that we need on the board."

AI might well follow the trajectory of data privacy in this regard, Chesnut said. The issue wasn't even a blip for most companies 20 years ago, but after a sharp rise in regulation, litigation, and public attention, most boards now have at least one member well-versed on privacy considerations.

"Every time a new area emerges that impacts how businesses operate and what they do, you need to ask the question of, how important is it that someone on the board understand this?" he said.

Cybersecurity governance can be a good model for boards in AI governance, GitHub's McKinley said, "as AI cybersecurity opportunities and threats should be top of mind."

And when CLOs talk to their boards about AI governance, they should frame it as "part of a broader, ongoing narrative about data governance," which includes cybersecurity and the EU's General Data Protection Regulation, McKinley said.

"By positioning AI governance in this way, CLOs can help boards understand that the principles they're already familiar with—transparency, accountability, and security—are simply evolving to meet new challenges."

Salesforce's Niles said boards would do well to devote a committee specifically to tech issues.

For the largest multinational companies and others that need to grapple with this broad set of issues, Niles said, "You're going to have to take a fresh look at how you want to organize the work of the board."

5 How can AI make my job easier?

Of course, AI isn't just a potential risk. It's supposed to make in-house legal teams more efficient in their own jobs.

Bloomberg Law Insight

Boards Need to Weigh Risk and Value With Generative AI Deployment



By Lara Abrash,
Deloitte

"In collaboration with management, boards should take up a strategic and intentional mindset, particularly as it relates to structure, skills development, and trust in technology."

[Read more.](#)

Clifford Chance's Kornbacher said clients often ask, "Can you tell us what you're doing and how you're using AI for your legal services?"

Microsoft's legal department is using its own and third-party tools for transactional processes, legal advisory services, and regulatory compliance and risk management, said Hossein Nowbar, the company's chief legal officer.

"For instance, our new internal Copilot Platform provides instant access to information across our data sources such as existing outside counsel (OC) memoranda to jumpstart our research, increase our agility, and reduce costs," he said in an email.

Gartner's in-house counsel clients are particularly interested in how generative AI can help them in contract management, Chris Audet, chief of research in Gartner's assurance practice, said in an email. Document summarization—including for litigation—is another promising use case for the near future, he said.

Gartner ranks AI's propensity for errors as among the highest risks for legal uses, along with risks related to data privacy and confidentiality.

But while hallucinations in legal AI tools are a well-publicized danger, Kornbacher said many lawyers looking

to invest in the new tech are more concerned about whether the tools actually live up to their promises.

There's a lot of "AI-washing," she said.

Most lawyers will be able to figure out if the tool hallucinates, she added.

"Rather than hallucination and discrimination, the risk is buying very expensive tools that don't actually increase efficiency or help productivity," she said. Instead they "take valuable time away from already lean and mean in-house teams."

There are no commonly accepted, standardized metrics for evaluating legal AI tools on qualities like accuracy and time saved.

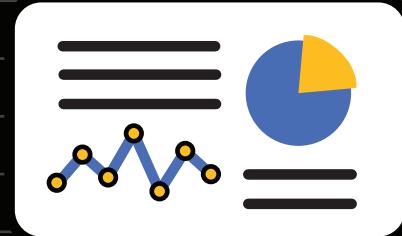
Microsoft's legal department assesses tools based on "effectiveness, including the tool's ability to drive cost savings, improve the speed and accuracy of legal tasks, and enhance overall productivity," Nowbar said. His team also looks for tools that can integrate with their existing systems.

Gartner is advising clients to identify the legal workflows where productivity gains of more than 10% are possible with the use of LLMs, Audet said.

OnPoint

Federal agencies are working to address issues from rapidly developing AI, a raft of proposals sit with Congress, and states are passing laws targeting risky AI uses. This slide deck sums up some of the key issues.

Download Now



Bloomberg Law

FEATURE STORY

Assigning Blame for AI Harms Raises Tricky Liability Questions

By Isabel Gottlieb

When an AI system goes awry and someone gets hurt, who's responsible? The company that made the technology, or the one that used it?

Courts, legislators, and government agencies are beginning to chip away at tricky questions about where liability lies when AI leads to harms, but the answers so far differ.

"Nobody knows" exactly where the liability falls, said Craig Shank, a consultant on AI ethics and formerly VP of corporate, external and legal affairs at Microsoft, a role he left in 2019. "There's not been enough litigation yet."

The EU may eventually legislate responsibility in a follow-on to its broad AI Act. In the U.S., it appears, the question will most likely be answered in federal agency regulation, in contracts between businesses, and in court.

The answer is murky because so many hands go into making an AI product that it's not clear how much influence each party has on the eventual outcome.

'Turtles all the way down'

If an AI system causes harm, "You're going to look backwards at who was

involved in creating the outcome of that system" said Lee Tiedrich, a professor at Duke University and an adviser to the OECD's AI work. "There are potentially a lot of different players in that ecosystem."

For example, the developer of an AI model may have purchased the data used to train the model. Then the company that purchases that AI model may fine-tune it with its own data. If there's a claim that the model produced discriminatory outcomes, it could be difficult to pinpoint whose data, training, or evaluation process is to blame.

"I'm not sure there are enough people with sophisticated enough knowledge of how AI works to be able to parse that—or be able to explain it to a jury even if they could parse it," said Thomas Magnani, a partner at Arnold & Porter. "So we're just in uncharted territory right now."

"The expression 'turtles all the way down' does come to mind at times when you look at all of this," Shank said. "But there are a lot of turtles in between me and the model."

Among the early court cases is a lawsuit alleging that automated tools from HR platform Workday Inc. discriminated against a job applicant. The litigation is ongoing, but there are early indications

Stay Up to Date Courts Orders



Some courts prohibit the use of generative AI while others direct attorneys to disclose and confirm the accuracy of any content that was generated by AI.

Track judicial standing orders.

Bloomberg
Law

NOW AVAILABLE
AI Legal Issues Toolkit

GO

that employers and vendors of HR AI tools could end up sharing liability when there are bias claims.

Regulators are studying the question, too. The Federal Trade Commission has asked in proposed rules whether AI platforms where deepfake video, audio, or images can be created should be liable if they know or have reason to know their services are being used to create deceptive content.

Policymakers who want to regulate AI liability have a range of choices, Tiedrich said. They could impose strict liability, though that would be a departure from where things stand today, she said. They could make platforms responsible for harms they should have known were possible, or raise questions about the companies' actual knowledge.

Rules like the FTC's proposed deepfake regulation could also turn to a notice-and-takedown model like the Digital Millennium Copyright Act, she said—but that also raises questions about false notices.

European Union policymakers are working on an AI liability directive that would clarify these questions—at least in the EU's 27 member states—but that's still a work in progress. The EU AI Act includes detailed compliance requirements for the developers and deployers of AI models that are at high risk of causing harms, but it doesn't clarify who would be sued when those harms occur.

Hashing it out in Contracts

Developers of general-purpose AI systems often disclaim all liability in commercial contracts and seek indemnification from users of their tools, Magnani said.

"The default for a long time is going to be, 'Buyer beware—caveat emptor,'" he said: The customer understands the risks of using the tool, and is using it to assist their work, not replacing their own judgment.

Copyright is an exception. After a spate of copyright infringement suits against generative AI tools, last year major AI players like Microsoft and Adobe said they'd indemnify their users for copyright infringement claims, but with conditions.

While AI providers initially sought to put the risk on the buyer, practices may shift in the coming years, said Daren Orzechowski, a partner at A&O Shearman.

"I think we're in a window of maybe two-to-three years where there's going to be a lot of churn within these deals—where, frankly, the terms that people are advocating for are all over the place," he said.

To help them with AI contracts, lawyers are looking to how similar questions were answered for recent technologies like cloud computing and blockchain, he said.

"We can't always wait and see," Orzechowski added. "We have to figure out how to allocate the risk now."



**Artificial Intelligence
Newsletter**

SIGN UP

Bloomberg Law | Bloomberg Tax | Bloomberg Government

PRACTITIONER INSIGHT

GCs Can Set Tone, Serve as Guides as Companies Navigate AI Era

By Sebastian Niles, *Salesforce*



The AI revolution has prompted a complex governance, technology, and regulatory thought exercise in how

organizations build trust around its use. Legal leaders and their teams have a key responsibility to set parameters around implementation and ensure all employees know how to interact with the technology—ethically and responsibly.

Customers and stakeholders are increasingly demanding transparency in AI systems, and general counsel can help businesses proactively meet and exceed those expectations. Here are some steps legal leaders can take to help their organizations deploy AI while building and maintaining stakeholder trust.

Establish Clear Guidelines and Governance

Absent any nationwide and global AI regulation, corporate legal teams must help their organizations navigate complex and uncharted territories. This starts with creating clear operational guidelines and instituting proper oversight.

These policies should account for several factors, including how the technology will enhance operations, privacy protocols—

especially if customer data will be involved—and potential associated risk factors.

To help apply and ensure adherence to these policies, proper oversight is key. Many organizations have implemented trusted AI councils with representation across the business—from legal to product experts—to help evaluate opportunities and use cases for implementing AI applications.

These experts should also be part of the product development cycle from the very beginning to help ensure effectiveness measures and ethical guardrails are part of AI system development from its inception—and not just an afterthought.

Embrace Reskilling and Upskilling

It is impossible to effectively set guidelines around the technology without a comprehensive understanding of how it works. Legal leadership should make an effort to source talent for their teams from various backgrounds, including people who have experience working with AI. They should also facilitate opportunities for their current team members to upskill in AI education.

Personally, I've taken time to dive into the intricacies of AI through a variety of avenues. Aside from participating

In Focus: Artificial Intelligence

NEW

Access news and resources across practice areas to track how lawmakers, courts, and agencies are addressing AI impacts.

GO

in our internal AI summits and product development, I've taken online trainings offered by our learning platform, attended informative conferences, and engaged with my peers and colleagues in in-depth discussions. I also encourage my team members to sign up for online or in-person courses or conventions they come across to help sharpen their own AI skills and learn how other organizations are approaching AI. But most importantly, I encourage our teams to share learnings, feedback, and insights with each other and our internal stakeholders so that we can help design the right AI future for the company and our customers.

Engage in a Multi-Stakeholder Approach

As AI continues to evolve, it is imperative that governments, businesses, and civil society organizations partner to build the right frameworks and guidelines.

General counsel and legal teams should partner with their government affairs colleagues to inform policymakers during the development and enforcement of generative AI policies and risk-based AI regulation.

At Salesforce, we regularly consult with both internal and external resources—including the National Institute of Standards and Technology's U.S. AI Safety Institute Consortium—to guide us in this important work.

We also regularly engage with policymakers across the globe on these crucial governance discussions, contributing to important global

gatherings such as the UN General Assembly and AI Safety & Innovation Summits in England and South Korea, and U.S. institutes.

Proactively Communicate with Customers

Legal leaders can play a critical role in communicating to customers how organizations are using AI. Companies should disclose generative AI use to their customers and should be prepared to explain how their AI systems make decisions—not only because of regulatory compliance but also because it's integral to building trust with customers and stakeholders.

Additionally, working with customers to provide clear, understandable explanations of AI decision-making processes is becoming increasingly important, especially in regulated industries.

At Salesforce, the Office of Ethical and Humane Use and legal teams collaborated to launch an AI acceptable use policy to provide greater clarity for our customers about how we use AI and promote safe and trusted experiences for all using our technologies.

We are at a critical turning point in the AI revolution, and, just as they have in areas of corporate compliance, general counsel have an opportunity to set the tone at their companies, lead by example for customers, and be the AI pioneers in their organizations.

"General counsel and legal teams should partner with their government affairs colleagues..."

In-House Counsel Newsletter



SIGN UP

PRACTITIONER INSIGHT

Open Dialogue Sets Up In-House Team for AI Adoption Success

By Amy Yeung, In-House Counsel Executive



In-house counsel voices have become an integral part of corporate conversations about implementing artificial intelligence. Those discussions range from learning about employees' individual experiences to addressing agenda topics in the boardroom.

As you use these insights to guide executive leadership on AI implementation, here are some operational considerations to build a strong foundation.

Feedback Loop

Collecting broad feedback across the organization can encourage buy-in and surface new use cases. Even where a team has previously identified multiple use cases, cultivating a precedent of continuous feedback encourages technology adoption, provides an opportunity for education and training, and develops a forward-looking innovation mindset.

It also opens the door to discuss the subtle but important distinction between replacing roles through technology and implementing technology to amplify individual capabilities and execute on outcomes. Collecting feedback can lead to discussions that keep messaging optimistic and goal-oriented, define expectations in a period of change, and promote a collaborative environment.

For example, I led an initiative implementing a contracts management and commercial templates system. The legal team set up a series of open forum meetings for department leaders and their invitees to share initial objectives

and receive feedback. In addition, we coordinated meetings with each of the leaders and asked them to designate a single point of contact who could communicate in-progress goals and provide ongoing feedback.

Once our primary objective was set, the process of vetting several vendors provided an opportunity to explore historical data trends and incorporate data analytics. This analysis involved our legal team, commercial teams, and vendor teams, and through it we learned that other groups were interested in tracking contract sections. The legal team came away with a greater understanding of department goals and further developed an opportunity to pursue additional long-term gains for the organization.

All this occurred before we had invested significant time and energy into the software integration, and receiving frequent feedback meant we could incorporate requests early into our project management scope and timeframe.

Pilot Program

Pilot programs are an early step in implementation, and in-house counsel can ensure the company gets the maximum benefit from them. Selecting simple use cases can promote additional feedback and provide more time to evaluate regulatory requirements.

Consider ways to use resources efficiently:

- Could a subset of documents be used to test the software, limiting enterprise involvement while still receiving enough data for meaningful feedback?

Stay Up to Date State Laws



Track state laws and regulations related to AI.

- Is there a small group who is well-suited to try out earlier versions of the software?
- Is there a portion of the software that could be implemented before intensive customization is required?

Consider another vantage point: de-risking scenarios for testing software and use cases.

- Could a neutral set of data be used to test the software?
- Could fake or anonymized data be used to test analytics?
- If anticipated data uses will require consumer disclosures, could software methodologies be validated by the engineering team before the disclosure changes are required?
- Do certain scenarios trigger prudential or regulatory review? If so, are there any scenarios that could be initiated in limited scope, before oversight is triggered, and still achieve employee and operational feedback?

As the legal and regulatory leader, take special care with certain AI applications as they may require greater compliance effort up front. For example, some AI technology includes the ability to collect data analytics for predictive purposes, which can provide powerful marketing insights targeting certain audiences. In this situation, a company is treating one group of individuals differently from another group, which may trigger additional legal, prudential, and regulatory requirements.

Crafting a pilot program that accomplishes the business value without

triggering these requirements will help avoid additional risk until it is truly necessary. The pilot program period also creates a defined space for all groups to proactively consider compliance requirements going forward using real-time insights from the pilot.

Regular Review

Technology changes fast. Today's technology setup won't and shouldn't look the same 12 months from now, and hopefully not six months from now either.

Plan future check-ins for usage and user experience. In my experience, a quick pulse survey is important to gauge if an initiative is generally aligned with expectations but doesn't do enough to understand how technology is impacting employee experiences. There should be continued discussion forums with the full leadership team reinforcing corporate priorities through this initiative.

Feedback also can alert in-house counsel on evolving compliance and regulatory needs. For example, a gradual shift in using different data sets may now require greater privacy disclosures and corresponding data governance adjustments. And the vendor's adjustments to the software may create new risks, requiring potential contractual changes. Ensure you have created a framework for ongoing reviews along with the compliance and audit teams.

Implementing AI is an exciting business opportunity and a chance for you as a leader to lean in as a key advisor and embrace a long-term perspective. With these strategies, the legal team will be a critical part of this AI paradigm shift.

Bloomberg Law Insight

Reskilling Is Required to Realize Generative AI's True Potential



By **Stephanie Corey** and **Brandi Pack**, **UpLevel Ops**

"Generative AI is transformative, but some early adopters are making a critical error: They lack a comprehensive reskilling strategy."

[Read more.](#)

Write for us!



[Click Here for Author Guidelines](#)

Bloomberg Law

PRACTITIONER INSIGHT

How In-House Legal Departments Should Engage with Generative AI

By Allen Waxman, Barclay T. Blair, Danny Tobey, **DLA Piper**



The promise of generative AI has brought a mix of anticipation and uncertainty to legal departments. And for good reason—it has potential to fully transform legal practice.

This technology holds promise for greater efficiency and cost savings, the ability to bring more legal work in-house, and reduced time spent on lower-value tasks. Yet the familiar litany of risks associated with its use include inaccuracies and hallucinations, data privacy and security, confidentiality, and IP concerns.

Still, converts appear to be growing rapidly. Pressure to adopt generative AI as a capability will no doubt continue to grow as the technology improves, as it gains greater acceptance in the marketplace, and as other departments (read: finance) seek the allure of cost savings.

A 2023 Accenture study reported that 100% of the work done by attorneys is exposed to transformation by generative AI, the only job category in the study to confront such significant change.

So how is the general counsel's office to respond? Deliberately. The very same care that distinguishes the legal department's success in so many areas should also distinguish the department's own engagement with generative AI.

However, moving deliberately doesn't have to mean moving glacially. After all, AI may not replace lawyers, but lawyers who use AI may replace those who don't.

We offer three factors legal departments ought to consider in adopting generative AI for their own practice.

What are you trying to accomplish by using generative AI?

And what pain points and opportunities are you trying to solve for? Answering these questions should be a team sport. Engage your group to identify the most meaningful use cases. As the tech world says: Love the problem, not the solution. In doing so, you are likely also to surface your team's concerns and uncertainty over using the technology.

Gathering input will not only lead to better thinking, including on how best to safeguard around risks, but also lead to greater buy-in for whatever use cases are pursued. In fact, some companies are going even further and surveying their outside counsel to leverage their experience with use cases as well.

In identifying use cases, it is important to balance the gains to be achieved against the effort to onboard, integrate, and validate the technology for your purposes.

What is the most appropriate technology to meet your needs?

Begin with an understanding of the technologies available. Determine what others are using within your industry, your outside law firms, and through consultancies.

Next, consider whether you want to work with generic models designed for all users or those intended for legal use cases. The latter includes existing legal tech providers incorporating generative AI into their platforms, legal information providers developing generative AI capabilities on top of their knowledge bases, and large law firms creating their own capabilities.

While the technology evolves, it may be worth engaging in a mix of approaches to hedge your risk. Ultimately, however, and as consulting firm Deloitte emphasized in its guide for engaging with generative AI in corporate legal departments, "The real potential lies in creating legal specific capabilities, which could be achieved by enhancing the data input into generic models, or by creating dedicated legal models." In our experience, we have seen this potential best realized by focusing multi-purpose models on highly targeted uses to maximize quality, consistency, and control.

And don't be afraid to experiment with whatever models you pursue. Engage in proofs of concept with different providers. In so doing, don't just focus on the accuracy and reliability of the tool tested, but understand the tool's data privacy and security implications.

Also, how does the technology treat the inputs you provide to it; be sure to understand whether it is safe to input your IP into the tool or whether you may unintentionally expose that IP for someone else's use.

Finally, assess whether the technology will be able to transition seamlessly into the workflows of your team. Innovative technology that doesn't integrate with the way your lawyers work is a cool demo that just never gets applied.

How should you onboard the technology you obtain?

Just as you may be doing for other departments in the company that use AI technology, you will want to set up a policy for who can use the tools as well as operating rules with any limitations on data usage. The operating rules should reflect considerations of the data privacy, security, and IP issues raised above, and the purposes of the tools. Ensure the policy is pragmatic so that it allows for innovation with appropriate guardrails.

You also will want to test the tool to validate its continuing accuracy as well as any biases that surface and monitor for any other concerns that might arise.

Training and educating your colleagues about the technology and the policy is important so they understand the best way to implement the tool and any limitations. Keep in mind the American Bar Association's recent pronouncement that lawyers should maintain "relevant technological competence, which requires an understanding of the evolving nature" of generative AI.

Finally, no one is better at anticipating that things will go sideways than the legal department. There may be cost overruns, delays in onboarding, snafus with the technology. The key to resilience is transparency with your legal department colleagues, with procurement, with finance, and with any other stakeholder in your process.

Keep them apprised of the rationale for your use cases (factor 1, above), the technology you are seeking (factor 2), and your intended onboarding process (factor 3). In this way, you will have less explaining to do when necessary and more investment in your intended outcomes.

Tools and Resources

Identify and mitigate risk, plan for the unexpected, and double check your work with these tools from Bloomberg Law's Practical Guidance. Click on any of the items below to read more.

- [Managing AI Compliance & Governance](#)
- [Strategies for Navigating AI Legal Risks](#)
- [M&A Pre-Transaction Review of AI](#)
- [Health AI Vendor Due Diligence](#)

PRACTITIONER INSIGHT

E-Discovery's Next Frontier Is Your AI Tool's Output, Prompt Log

By Thomas Barce, **FTI Technology**



It's happening. Leading organizations are forming committees, budgeting resources, performing research, conducting pilots, and beginning early implementations for generative artificial intelligence.

Many legal teams are under pressure from leadership to be among the early adopters amid efforts to gain value from these advancements. As new applications enter the market, there's an expectation that teams will quickly onboard tools, and in turn, realize measurable gains.

But there are also concerns about the risks, and legal teams are torn between pursuing innovation and maintaining a strong risk management posture.

Generative AI's entrance into enterprise environments has created a new dimension of company information and potential liability that many organizations aren't quite sure how to handle. Information governance controls are now required for an uncharted category of records—namely “interactions,” which are logs of prompts used to query AI tools. New discovery rules and processes must be established for data categories that haven't been discoverable, including interactions and company documents created entirely by a machine.

Legal teams must now address governance and compliance, as well as e-discovery readiness, when implementing generative AI. They also must proactively map out policies for how employees may interact with these tools and where the underlying interactions are stored. This way, they can be properly retained, monitored for compliance purposes, defensibly disposed, and

preserved as needed for future legal discovery requirements.

While organizations are working toward generative AI use cases, more risk-oriented controls need to be considered in parallel:

- What happens to generative AI interactions within the organization?
- Where are interactions stored?
- Are they being retained or disposed of?
- Are they being monitored?

Organizations should monitor and manage the use of AI tools in the same way that email, chats, and other established communications require oversight to catch misuse or violation of varying regulatory obligations. Legal teams can:

- Conduct information security and third-party risk management audits to confirm the extent that AI tools comply with company policies and regulatory requirements for other technologies, systems, and providers.
- Establish contractual controls—such as indemnity clauses or stipulating that no client, confidential or sensitive data may be used in large language model training—to provide general protections while teams sort through unknowns within the technology.
- Implement abuse monitoring capabilities to notify compliance teams if employees are making suspicious or inappropriate queries

"Generative AI's entrance into enterprise environments has created a new dimension of company information and potential liability..."

of the AI tool while balancing retention, protection, and access of data that may be monitored—especially including proprietary or personal information.

- Review access controls for company records. Many off-the-shelf AI tools provide users with widespread and unfettered access to query company documents and information. Every application in the system must be checked for access controls to be sure that queries of the AI system won't result in users viewing information they shouldn't be able to access.
- Label policies for documents and information categories that must be treated with varying levels of confidentiality or protection.

Employee interactions with and output from generative AI tools are creating a data set that could come into scope in e-discovery. For example, generative AI tools store artifacts that may introduce e-discovery implications when data related to or from the tools intersects with a dispute or investigation.

When prompts and interactions are stored, they become a potential company document, communication, or record that could be considered as evidence in litigation or a regulatory investigation. The number of tools creating these new or unknown data artifacts, with limited visibility or accessibility, is a potential Pandora's box.

Just like with other forms of emerging data, it will be challenging to preserve these artifacts, defensibly collect them, process them into an e-discovery tool, and render them useful for analysis and review. It also will require technical expertise. Legal teams may also need to irrefutably distinguish between human-generated content and AI-gener-

ated content when entering these items into evidence in a legal matter.

Legal arguments about interactions with generative AI, their qualification as communications, and whether they may be subject to discovery will be complicated, if not contentious. Regulators and litigators quickly identified modern forms of communication such as Slack, Zoom, Teams, various mobile messaging apps, and other ephemeral messaging platforms as relevant sources of evidence. It is also plausible that as generative AI becomes more mainstream, data from these systems will likewise be sought to support fact finding throughout the course of discovery.

Many organizations are currently at varying phases of generative AI implementation, with most at either the proof of concept or pilot phase. Functionality and controls change constantly, so organizations need their AI governance and e-discovery readiness programs to be built for adaptability and continual testing.

Until we reach a high degree of certainty with this technology, generative AI tools likely are or will create data that's not being properly controlled. Organizations must activate the appropriate analyses and put the right controls in place. Meanwhile, it's just as probable that some of those controls are yet to be conceived.

Considering the typical 12- to 24-month lag time between the adoption of emerging technology and its appearance in relevant litigation or investigations, the clock is ticking. What seems brand new today may be the linchpin in tomorrow's court case.

Whether an organization currently sanctions the use of generative AI for business, potentially relevant evidence from generative AI has probably already been created.

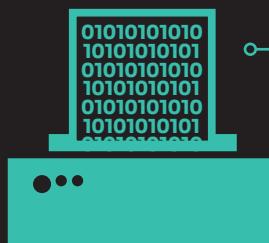
"The number of tools creating these new or unknown data artifacts, with limited visibility or accessibility, is a potential Pandora's box."

Glossary of Terms

1·2·3

Algorithm

A step-by-step process used to solve a problem.



Machine Learning

The process of feeding data into computer algorithms so they get more refined and sophisticated over time.

Natural Language Processing

The branch of AI that helps computers to understand, process, and generate speech and text the way a human would.



X	X	X
✓	X	X
X	X	✓

Deep Learning

The most common form of AI, in which software is taught to classify something such as a video or a loan application from a very large set of labeled data.

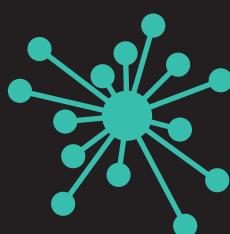
Chatbots

These are products that can hold advanced, human-like conversations with people about anything from historical trivia to lists of creative recipes using a watermelon.



Large Language Models

The backbone of natural language processing that can summarize and generate text using information from all over the internet. Perhaps the most well-known is OpenAI's GPT-4.



Generative AI



This refers to the production of entirely new creative works—pictures, music, text, poetry—from simple prompts after AI is trained on vast quantities of pre-existing material.

Hallucinations

The phenomenon by which AI chatbots may confidently provide false information (sometimes ludicrously so) in response to a prompt.



*Adapted from "A Cheat Sheet to AI Buzzwords and Their Meanings: QuickTake," Bloomberg News.

Bloomberg Law

Advise your business with confidence

Bloomberg Law is the only platform that empowers your legal team to expand their expertise, tackle unfamiliar issues in-house, and respond to stakeholders faster.

Trusted news and analysis

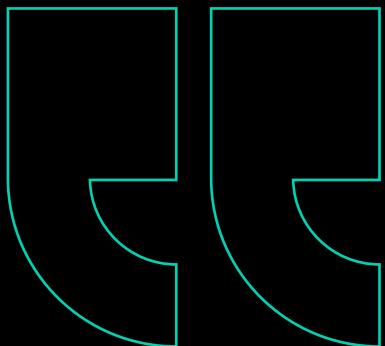
Primary and secondary sources

Step-by-step expert legal guidance

Comprehensive practice area coverage

Time-saving workflow tools

[Request a demo](#)



Bloomberg Law allows us to do much of the work in-house. It gives our department the confidence that we have the correct information. It provides quick answers that save us from having to call outside counsel. As a result, we were able to decrease our annual outside counsel spend by 15-20%.

Deputy General Counsel // Woodside Group, LLC



About Bloomberg Law

Bloomberg Law helps legal professionals provide counsel with access to action-oriented legal intelligence in a business context. Bloomberg Law delivers a unique combination of Practical Guidance, comprehensive primary and secondary source material, trusted content from Bloomberg Industry Group, news, time saving practice tools, market data, and business intelligence.

For more information, visit pro.bloomberglaw.com.

Bloomberg Law